

Implementing Dual Pivot Quicksort in C++23

Validating Theoretically Superior Sorting Strategies in Modern Systems

Student: LIN Zhanzhi

Supervisor: Dr. CAO Yixin

Project: Capstone Interim Assessment 2025/26

Presentation Agenda

- **Background & Motivation:** Why Dual-Pivot?
- **Sequential Performance:** Beating std::sort by 15-28x.
- **Parallel Architecture:** Work-Stealing & The "Memory Wall".
- **Future Roadmap:** Breaking the wall with AVX-512.

The Status Quo vs. The Innovation

Standard C++ (Introsort)	Our Project (Dual-Pivot)
Single Pivot	Dual Pivots (P_1, P_2)
2 Partitions ($< P, \geq P$)	3 Partitions ($< P_1, P_1..P_2, > P_2$)
Optimized for CPU Cycles	Optimized for Memory Bandwidth

"Scanned Elements Model":

Moving elements is expensive (Memory Write). Checking them is cheap (CPU Read).

3-Way Partitioning = Fewer Swaps = Better Cache Efficiency.

Project Objectives

1. Modernize

- Create a **Generic C++23 Library**.
- Use Concepts (`std::sortable`) for type safety.

2. Accelerate


- Outperform `std::sort` on sequential benchmarks.
- Achieve scalable parallel performance.

3. Analyze

- Identify hardware limits.
- Investigate the "**Memory Wall**" in parallel sorting.

Baseline Results: Random Data

- **Metric:** 64-bit Integers, $N = 10^3$ to 10^7 .
- **Result:** Consistent 10-15% Speedup vs `std::sort`.



DPQS (Green) < `std::sort` (Red)



The "Run Merger" Optimization

Why sort what is already sorted?

- **Strategy:** Adaptive Run Merging (TimSort-style).
- **Structured Data Performance (10M elements):**


Distribution	DPQS Time	std::sort Time	Speedup
Organ Pipe	22.76870 ms	654.35300 ms	28.7x
Sawtooth	28.27810 ms	239.18200 ms	8.5x

(Data for int, Size: 10,000,000)




Robustness: The "Fat Partition" Problem

- **Challenge:** Duplicate pivots usually degrade Quicksort to $O(N^2)$.
- **Solution:** 3-Way Partitioning clusters duplicates:
 - Region 2 ($P_1 \leq x \leq P_2$) naturally absorbs equal keys.
- **Result:** Performance is **invariant** across 10%, 50%, or 90% duplicates.




Parallel Architecture: V3 Work-Stealing



Strong Scaling Results

Setup: $N = 10,000,000$ Integers.

Threads	Speedup (vs Sequential)	Efficiency	Phase
2	2.03x	101%	Super-Linear
4	3.52x	88%	Linear
8	4.98x	62%	Diminishing
16	5.24x	33%	Saturation



The Bottleneck: Hitting the "Memory Wall"

Why does speedup plateau at 5.2x?

- **It's NOT** Software Overhead (Locks are minimal).
- **It IS** Memory Bandwidth Saturation.
 - Sorting is $O(N)$ Read/Write intensive.
 - 16 Threads starve the memory controller.
 - *Adding more CPU cores cannot fix a data supply shortage.*

Engineering Quality: C++23 Standard

We rely on **Modern C++ Concepts** to ensure type safety and API compatibility.

```
// include/dual_pivot_quicksort.hpp

template<std::random_access_iterator RandomAccessIterator>
void dual_pivot_quicksort(RandomAccessIterator first, RandomAccessIterator last) {
    // Validated at compile-time by C++20 Concepts

    if (first >= last) return;
    // ...
```

- **Generic:** Works with `std::vector`, `std::array`, `std::deque`.
- **Safe:** 64-bit support preventing integer overflow.

Roadmap: Semester 2

Goal: Break the Memory Wall.

1. ⚡ SIMD Vectorization (AVX-512)

- Use **Non-Temporal Stores** to bypass cache and write directly to RAM.
- Expected to double effective bandwidth.

2. 🧠 Memory-Aware Scheduling

- Prioritize task stealing based on memory locality (NUMA awareness).

3. 📈 Advanced Benchmarking

- Compare against state-of-the-art **PDQSort**.

Conclusion

1. Sequential Success:

- 15% faster on Random Data.
- 28x faster on Structured Data.

2. Parallel Foundation:

- 5.24x Scalability achieved.

3. Rigorous Analysis:

- Identified **Memory Bandwidth** as the true limit.

Next Step: Hardware-native optimizations to push beyond the wall.

Thank You.