

# Meeting Note 4

---

**Date:** 2026-01-02 **Time:** 14:00 - 15:00 **Location:** Supervisor's Office (PQ706) **Attendees:** LIN Zhanzhi, CAO Yixin (Supervisor)

## 1. Objectives

- Provide a comprehensive progress report on the completed sequential and parallel modules.
- Present the benchmarking results showing performance superiority over `std::sort`.
- Validate the scope and structure for the upcoming Interim Report.
- Discuss technical decisions regarding SIMD and parallel architecture.

## 2. Progress Report

- **Algorithmic Core:**
  - Successfully implemented **3-way partitioning** with "Median-of-5" pivot selection, achieving ~20% fewer swaps.
  - Integrated **Adaptive Run Merging** (TimSort-style) to handle pre-sorted data efficiently.
  - Implemented **Introsort Strategy** (fallback to HeapSort) to guarantee  $\$O(N \log N)\$$  worst-case safety.
  - Added **Counting Sort** optimization for small integer types (`byte`, `short`).
- **Parallel Architecture:**
  - Deployed a **Work-Stealing Thread Pool** for dynamic load balancing.
  - Completed contention analysis and mutex optimization.
- **Robustness:**
  - Ensured namespace safety and fixed integer overflow issues for large arrays.
  - Verified generic support with custom comparators.
- **Infrastructure:**
  - Established a custom **Benchmarking Harness** with "Destructive Testing" support.
  - Performed rigorous sample size analysis to ensure statistical significance.

## 3. Discussion & Feedback

### Key Discussion Points

- **Sequential Performance:** Reviewed the data showing DPQS outperforming `std::sort` on random integers and achieving massive speedups on structured data (e.g., Organ Pipe). Supervisor acknowledged the efficacy of the "Run Merging" strategy.
- **Parallel Scaling:** Discussed the "Strong Scaling" results. The linear scaling up to 4 threads and saturation beyond 8 threads (due to Memory Bandwidth) was accepted as a valid finding for the Interim Report.
- **SIMD Feasibility:** Presented the decision to prioritize Parallelism over SIMD for Semester 1. Supervisor agreed that SIMD complexity is high and fitting as a "stretch goal" for Semester 2.
- **Comparison Targets:** Confirmed that comparing against `std::sort` and `std::stable_sort` is sufficient for the interim phase. Future work can include comparisons with TBB or OpenMP.

## Supervisor Feedback

"The algorithmic tuning, especially the Adaptive Run Merging, is a strong point. Make sure to highlight *why* 3-way partitioning helps with duplicates in your report."

"Your identification of the 'Memory Wall' in the parallel section is critical. Ensure your Interim Report explains this hardware bottleneck clearly—don't just show the graph, explain the saturation."

"The scope for the Interim Report is appropriate. The focus on 'Implementation & Validation' is solid. You don't need AVX-512 yet; getting the thread pool working correctly was the right priority."

## 4. Issues & Challenges

- **Parallel Overhead:** Discussed the challenge of managing thread overhead for small-to-medium arrays.
  - *Resolution:* Masked by the dynamic threshold ( $N < 4096$ ) falls back to sequential), which effectively amortizes the cost.
- **SIMD Complexity:** Acknowledged that AVX-512 implementation is complex.
  - *Resolution:* Formally moved to Semester 2 plan as "Advanced Optimization".

## 5. Action Items

- Finalize the Interim Report based on the approved outline. (Due: 2026-01-09)
- Prepare the Interim Presentation slides, focusing on the "Speedup Analysis" graphs. (Due: 2026-01-09)
- Refine the work-stealing deque implementation for better load balancing in edge cases. (Due: Feb 2026)
- Begin preliminary investigation into AVX-512 intrinsics for the "Future Work" section. (Due: Feb 2026)

## 6. Next Meeting Plan

- **Target Date:** 2026-02-27 (Semester 2 Start)
- **Focus:** Detailed plan for SIMD implementation and advanced parallel optimizations.