

Meeting Note 1

Date: 2025-10-01 **Time:** 14:00 - 14:45 **Location:** Supervisor's Office (PQ706) **Attendees:** LIN Zhanzhi, CAO Yixin (Supervisor)

1. Objectives

- Review the initial draft of the Capstone Project Proposal.
- Discuss the feasibility and scope of implementing Dual-Pivot Quicksort in C++.
- Identify key academic and technical requirements for the project.

2. Progress Report

- Selected project topic: "Implementing Dual Pivot Quicksort in C++".
- Drafted a preliminary proposal outlining the basic implementation plan.
- Conducted initial reading on Yaroslavskiy's 2009 paper.

3. Discussion & Feedback

Key Discussion Points

- **Scope Definition:** The initial proposal focused primarily on the coding aspect. We discussed the need to elevate this to an academic research project by adding rigorous performance analysis and theoretical backing.
- **Language Standards:** Discussed the choice of C++23. Supervisor agreed it is a good choice to demonstrate modern software engineering practices (e.g., using `std::ranges` and concepts).
- **Benchmarking:** Simply implementing the algorithm is not enough. The core value lies in comparing it against existing giants like `std::sort` (Introsort) and PDQSort.

Supervisor Feedback

"The current draft is too engineering-focused. You need to justify *why* this is a research gap. Java adopted Dual-Pivot years ago; why hasn't C++? You should investigate the 'scanned elements model' by Wild to explain cache behavior, rather than just counting comparisons."

"Expand the methodology to include a systematic comparative analysis. Don't just measure time; measure cache misses and branch mispredictions."

4. Issues & Challenges

- **Theoretical Depth:** I admitted I was unfamiliar with the mathematical analysis of Dual-Pivot Quicksort (e.g., Aumüller's framework).
 - *Resolution:* Supervisor provided references to key papers by Aumüller and Dietzfelbinger to strengthen the literature review.
- **Benchmarking Fairness:** How to ensure a fair comparison between my implementation and the highly optimized STL?

- *Resolution:* We agreed to use a standardized benchmarking framework (like Google Benchmark) and test across various distributions (random, sorted, reverse).

5. Action Items

- Read Aumüller and Dietzfelbinger's paper on Dual-Pivot analysis. (Due: 2025-10-10)
- Read Wild's paper on the "scanned elements model". (Due: 2025-10-10)
- Rewrite the "Background" section to highlight the gap in C++ standard libraries. (Due: 2025-10-15)
- Expand "Methodology" to include specific metrics (L1/L2 cache misses, branch prediction). (Due: 2025-10-15)
- Submit the revised, final proposal. (Due: 2025-10-24)

6. Next Meeting Plan

- **Target Date:** 2025-11-05
- **Focus:** Review of the revised proposal and initial code structure.

Meeting Note 2

Date: 2025-11-05 **Time:** 14:00 - 14:30 **Location:** Supervisor's Office (PQ706) **Attendees:** LIN Zhanzhi, CAO Yixin (Supervisor)

1. Objectives

- Review progress on the initial C++ implementation of Dual-Pivot Quicksort.
- Discuss findings from the literature review (Wild and Aumüller papers).
- Plan the structure for the upcoming Interim Report.

2. Progress Report

- **Proposal:** Final project proposal submitted on Oct 24.
- **Literature Review:** Completed reading of "Dual-Pivot Quicksort: Optimality and Analysis" (Aumüller) and Wild's work on scanned elements.
- **Implementation:** Created a basic working prototype of 1-pivot and 2-pivot quicksort in C++.
 - Current status: Works for `std::vector<int>`.

3. Discussion & Feedback

Key Discussion Points

- **Code Quality:** Showed the initial code. Supervisor noted that while it works for `int`, it needs to be generic.
- **C++23 Concepts:** Discussed using `std::sortable` and `std::random_access_iterator` concepts to ensure type safety.
- **Interim Report:** Supervisor outlined the expectations for the Interim Report.

Supervisor Feedback

"Your current implementation is too specific to integers. You must use C++ templates and Concepts to make it a true library candidate. Ensure it handles custom comparators correctly."

"For the Interim Report, don't just paste code. Focus on the *design decisions*. Why did you choose this specific partitioning scheme? Use diagrams to explain the 3-way partition."

4. Issues & Challenges

- **Generic Implementation:** I am struggling slightly with the syntax for C++20/23 Concepts for custom iterators.
 - *Resolution:* Supervisor suggested looking at the `std::ranges::sort` implementation in the GCC or LLVM open-source repositories for reference.
- **Benchmarking Setup:** Setting up Google Benchmark is proving tricky with the custom build system.
 - *Resolution:* Advised to stick to a simple `CMake` setup to manage dependencies.

5. Action Items

- Refactor code to use C++ Templates and Concepts (Generic programming). (Due: 2025-11-20)
- Implement a "Benchmark Runner" using Google Benchmark. (Due: 2025-11-25)
- Draft the "Methodology" and "Literature Review" sections for the Interim Report. (Due: 2025-12-01)

6. Next Meeting Plan

- **Target Date:** 2025-12-03
- **Focus:** Review of Interim Report Draft and preliminary benchmark results.

Meeting Note 3

Date: 2025-12-03 **Time:** 14:00 - 14:45 **Location:** Supervisor's Office (PQ706) **Attendees:** LIN Zhanzhi, CAO Yixin (Supervisor)

1. Objectives

- Review the draft of the Interim Report.
- Analyze preliminary benchmarking results.
- Discuss requirements for the Interim Presentation Video.

2. Progress Report

- **Interim Report:** Completed 80% of the report (Introduction, Literature Review, Methodology).
- **Implementation:** Refactored code to be fully generic using C++20 Concepts. Added `BenchmarkRunner` class.
- **Results:** Initial benchmarks show DPQS is faster than `std::sort` for large arrays (>1M elements) of random integers.

3. Discussion & Feedback

Key Discussion Points

- **Report Structure:** Supervisor reviewed the draft. Pointed out that the "Methodology" section was too descriptive of the code. It needs to explain the *algorithm logic* more clearly (pseudocode/flowcharts).
- **Benchmark Noise:** The current results have high variance.
 - *Advice:* Run benchmarks on a dedicated machine or ensure background processes are minimized. Use "median" instead of "average" to filter outliers.
- **Future Work:** Discussed the plan for Semester 2. Supervisor strongly suggested exploring **Parallelization** (using `std::thread` or OpenMP) as the next major step.

Supervisor Feedback

"The results look promising, but a table of numbers is hard to read. Generate line graphs showing 'Time vs Input Size'. Also, include a 'Speedup' graph relative to `std::sort`."

"For the presentation video, make sure you *show* the sorting happening (maybe a small visualization or live demo of the benchmark running). Don't just read slides."

4. Issues & Challenges

- **Small Array Performance:** DPQS is slower than `std::sort` for small arrays (<100 elements).
 - *Resolution:* This is expected. Discussed implementing a hybrid approach (switching to Insertion Sort for small partitions), similar to Introsort. This will be a key task for the next phase.

5. Action Items

- Generate graphs for the Interim Report results section. (Due: 2025-12-10)

- Record and edit the Interim Presentation Video. (Due: 2025-12-15)
- Finalize and submit the Interim Report. (Due: 2025-12-18)
- Begin research on Parallel Quicksort strategies. (Due: Jan 2026)

6. Next Meeting Plan

- **Target Date:** 2026-01-15 (Semester 2)
- **Focus:** Parallel implementation plan and feedback on Interim Assessment.

Meeting Note 4

Date: 2026-01-02 **Time:** 14:00 - 15:00 **Location:** Supervisor's Office (PQ706) **Attendees:** LIN Zhanzhi, CAO Yixin (Supervisor)

1. Objectives

- Provide a comprehensive progress report on the completed sequential and parallel modules.
- Present the benchmarking results showing performance superiority over `std::sort`.
- Validate the scope and structure for the upcoming Interim Report.
- Discuss technical decisions regarding SIMD and parallel architecture.

2. Progress Report

- **Algorithmic Core:**
 - Successfully implemented **3-way partitioning** with "Median-of-5" pivot selection, achieving ~20% fewer swaps.
 - Integrated **Adaptive Run Merging** (TimSort-style) to handle pre-sorted data efficiently.
 - Implemented **Introsort Strategy** (fallback to HeapSort) to guarantee $\$O(N \log N)\$$ worst-case safety.
 - Added **Counting Sort** optimization for small integer types (`byte`, `short`).
- **Parallel Architecture:**
 - Deployed a **Work-Stealing Thread Pool** for dynamic load balancing.
 - Completed contention analysis and mutex optimization.
- **Robustness:**
 - Ensured namespace safety and fixed integer overflow issues for large arrays.
 - Verified generic support with custom comparators.
- **Infrastructure:**
 - Established a custom **Benchmarking Harness** with "Destructive Testing" support.
 - Performed rigorous sample size analysis to ensure statistical significance.

3. Discussion & Feedback

Key Discussion Points

- **Sequential Performance:** Reviewed the data showing DPQS outperforming `std::sort` on random integers and achieving massive speedups on structured data (e.g., Organ Pipe). Supervisor acknowledged the efficacy of the "Run Merging" strategy.
- **Parallel Scaling:** Discussed the "Strong Scaling" results. The linear scaling up to 4 threads and saturation beyond 8 threads (due to Memory Bandwidth) was accepted as a valid finding for the Interim Report.
- **SIMD Feasibility:** Presented the decision to prioritize Parallelism over SIMD for Semester 1. Supervisor agreed that SIMD complexity is high and fitting as a "stretch goal" for Semester 2.
- **Comparison Targets:** Confirmed that comparing against `std::sort` and `std::stable_sort` is sufficient for the interim phase. Future work can include comparisons with TBB or OpenMP.

Supervisor Feedback

"The algorithmic tuning, especially the Adaptive Run Merging, is a strong point. Make sure to highlight *why* 3-way partitioning helps with duplicates in your report."

"Your identification of the 'Memory Wall' in the parallel section is critical. Ensure your Interim Report explains this hardware bottleneck clearly—don't just show the graph, explain the saturation."

"The scope for the Interim Report is appropriate. The focus on 'Implementation & Validation' is solid. You don't need AVX-512 yet; getting the thread pool working correctly was the right priority."

4. Issues & Challenges

- **Parallel Overhead:** Discussed the challenge of managing thread overhead for small-to-medium arrays.
 - *Resolution:* Masked by the dynamic threshold ($N < 4096$) falls back to sequential), which effectively amortizes the cost.
- **SIMD Complexity:** Acknowledged that AVX-512 implementation is complex.
 - *Resolution:* Formally moved to Semester 2 plan as "Advanced Optimization".

5. Action Items

- Finalize the Interim Report based on the approved outline. (Due: 2026-01-09)
- Prepare the Interim Presentation slides, focusing on the "Speedup Analysis" graphs. (Due: 2026-01-09)
- Refine the work-stealing deque implementation for better load balancing in edge cases. (Due: Feb 2026)
- Begin preliminary investigation into AVX-512 intrinsics for the "Future Work" section. (Due: Feb 2026)

6. Next Meeting Plan

- **Target Date:** 2026-02-27 (Semester 2 Start)
- **Focus:** Detailed plan for SIMD implementation and advanced parallel optimizations.